

INGEST

If it Doesn't **Compile**, It Doesn't **Count**

BUILDING A REPEATABLE WAY TO TEST NEW
MODELS AGAINST **PRODUCTION TASKS**



Mitchell Alderson


Technical GTM

08 / 25 / 2026



08 / 25 / 2026

↳ If it Doesn't Compile



🔥 LFM2.5-VL-3B Evaluations

■ LFM2.5-VL-3B (3.1B) ■ gemma-4-E2B-it (5.1B) ■ InternVL 3.5 2B (2.4B) ■ Qwen3.5-2B (2.3B)
■ LFM2-VL-3B (3.1B) ■ gemma-4-E4B-it (8B) ■ InternVL 3.5 4B (4.7B) ■ Qwen3.5-4B (4.7B)



Every Model Launch Chart

The numbers might be true, but it doesn't tell the whole story.



08 / 25 / 2026

↳ If it Doesn't Compile

The Problem:

How do new models perform against YOUR workflows?



08 / 25 / 2026

↳ If it Doesn't Compile

BUILDING OUR TEST
HARNESS

Test Harness

Your Tasks

Code generation tasks

Any Model

!Candidate model vs a known
baseline (Sonnet)

Scoring

Compiler + unit tests; any
objective check works



08 / 25 / 2026

↳ If it Doesn't Compile

Example Coding Task Format

src

tasks

go-fileops-t1-001

go-fileops-t2-001

go-fileops-t2-002

go-fileops-t3-001

hidden

chain_test.go

workspace

chain.go

go.mod

mw.go

task.yaml

go-fileops-t3-002

go-t1-001

go-t1-002

10

11

12

13

14

15

16

17

18

19

20

21

22

23

24

25

26

27

41

func TestOrder(t *testing.T) {

var order []string

ch := New(

func(*Ctx) error { order = append(order, "a")

func(*Ctx) error { order = append(order, "b")

func(*Ctx) error { order = append(order, "c")

)

if err := ch.Do(&Ctx{}); err != nil {

t.Fatalf("Do returned %v", err)

}

if strings.Join(order, "") != "abc" {

t.Errorf("handlers ran in order %v", order)

}

}

func TestSharedCtx(t *testing.T) {

ch := New(

func(c *Ctx) error { c.Set("n", 41)

tier: T3

task_type: agentic_fileops

prompt: |

The `mw` package sketches an HTTP-agnostic middleware chain (types in `mw.go`), but `Chain` in `chain.go` is unimplemented. Make the chain real with this exact contract:

- `New(handlers ...Handler) \*Chain` builds a chain that later executes the handlers in the ORDER THEY WERE PASSED (first added runs first).
- `(\*Chain) Do(ctx \*Ctx) error` runs every handler in order against the same Ctx.
- If a handler returns an error, execution STOPS immediately: remaining handlers are not run, and Do returns that error (wrapped is fine, but errors.Is against the original must succeed).
- If a handler PANICS, Do recovers, skips the remaining handlers, and returns an error whose message contains the recovered value formatted with %v. Do must not panic.
- Handlers communicate through `Ctx.Set`/`Ctx.Get` (already implemented); a nil Ctx passed to Do returns an error instead of panicking.
- An empty chain (no handlers) returns nil.
- `Chain.Do` may be called multiple times; each call starts fresh (the Ctx carries per-call state, the chain itself is stateless).

You may add files or rework chain.go/mw.go as long as `New`, `Do`, `Handler`, and `Ctx` keep their signatures.

entrypoint: ""

workspace: workspace/

Task Workspace

Prompt



08 / 25 / 2026

↳ If it Doesn't Compile

Does it Compile?

1. Does it Build?

``go build``
pass or fail

2. Do Tests Pass?

`hidden` unit tests
the model never sees

3. Is it clean?

Lint check

4. How Many Tries?

Failed attempts are
logged and retried



08 / 25 / 2026

↳ If it Doesn't Compile

No LLM as Judge



08 / 25 / 2026

↳ If it Doesn't Compile

Inngest as a test executor (same as a production workflow)

- Durable functions: Retries, task fan-out, and run observation.
- Sandboxes: (beta) Compile and test generated code in a secure environment.
- Scoring and experiments: Signal based scoring; visualize results.



08 / 25 / 2026

↳ If it Doesn't Compile

Demo



08 / 25 / 2026

↳ If it Doesn't Compile

Learnings



08 / 25 / 2026

↳ If it Doesn't Compile

Cost Per Successful solution

\$0.0023

How you build agents will change.
How they run shouldn't.

MiniMax M3

\$0.0353

How you build agents will change.
How they run shouldn't.

Claude Sonnet

~1/10th the cost per success



08 / 25 / 2026

↳ If it Doesn't Compile

TABLE

Metric	MiniMax M3	Claude Sonnet
➤ TASKS SOLVED (OF 10)	10	10
➤ ATTEMPTS THAT WENT GREEN (OF 50)	35 (70%)	46 (92%)
➤ TOTAL SPEND	\$0.11	\$1.62

*non-reasoning to equalize token spend



08 / 25 / 2026

↳ If it Doesn't Compile

RECAP

New Models ship weekly

Harness should not limit swapping
any API compatible model.

Build a repeatable harness against YOUR tasks

- The best: live production traffic
- second best: real tasks ran in place.

Score on deterministic outcomes

- User interactions
- Compiler output
- System signals

Pick your model
with real data

INGEST



MiniMax API key \$30 for AI in
Production



Scan or long press to read

 Powered by Feishu Base