

INGEST

Ship Your First Eval

YOU'RE ALREADY RUNNING AGENTS IN PROD



Mitchell Alderson

Technical GTM

08 / 12 / 2026

INGEST

Who am I?

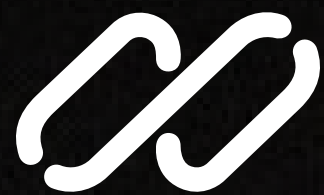
- 15+ yr software engineer
- Background building durable workflows by hand
- Filmmaker and Content Creator



Mitchell Alderson

Technical GTM

08 / 12 / 2026



08 / 12 / 2026

↳ Ship Your First Eval

Who is This for?

Inngest User

Or familiar with durable execution and the Inngest SDK

Comfortable with Typescript

Examples are in Typescript, but you can ask Claude!

Comfortable with Typescript

Examples are in Typescript, but you can ask Claude!




↳ Ship Your First Eval

Full Code Example on GitHub

Phase 1: starter agent in Inngest

Phase 2: First eval

Phase 3: advanced example





08 / 12 / 2026

↳ Ship Your First Eval

SHORT AGENDA

- What Are Evals?
- Picking Your First Eval
- Adding Scoring code to our Agent
- Viewing our Outcomes
- Where to Next?

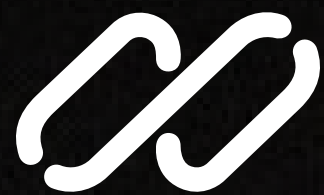
Pick First Eval

Adding Scoring Code

Viewing Eval Data

What is an Evaluation?

A single metric used to
measure agent output on a
specific aspect.



08 / 12 / 2026

↳ Ship Your First Eval

TYPES OF EVALS

Ideas for Initial
Metrics

LLM-as-Judge

Use another LLM to grade tone, helpfulness, or other subjective traits.

Tool Calls

Measure the number or frequency of tool calls the agent made.

Instruction Following

Compare the user input to the output, did the agent follow instructions?

User Intent

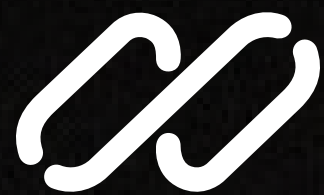
Capture user actions to grade the agent directly on human-in-the-loop feedback.

Structured Output Parsing

Parse the agent output, does it format to clean JSON?

Cost Optimization

Count token and word output, was the agent frugal with its output?



08 / 12 / 2026

↳ Ship Your First Eval

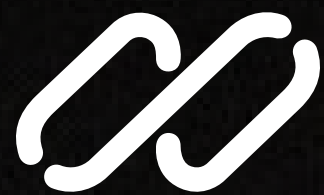
ONLINE VS. OFFLINE EVALS

Online

- Monitors and scores on live traffic as it happens
- Allows for experiments on real traffic
- With orchestration, allows for capture of human feedback
- Does not require an existing dataset of interactions

Offline

- Allows for regression and pre-deploy testing
- Does not affect any live traffic
- Requires dataset of “Golden Answers”
- Live edge cases may be missed



08 / 12 / 2026

↳ Ship Your First Eval

Why Inngest?

🕒 EVALS & OBSERVABILITY

Data collection, observability, and scoring

🔗 BUSINESS LOGIC

Agent code and business specific functions

🏗️ EXECUTION LAYER

Inngest Orchestration

🕒 RUNTIME

Node.js, Python, native binary, etc.



08 / 12 / 2026

↳ Ship Your First Eval

Deferred Scoring

Pick First Eval

Adding Scoring Code

Viewing Eval Data





08 / 12 / 2026

↳ Ship Your First Eval

Testing our Agent

Clicking the 👍/👎 will fire off the `chat/feedback.clicked` Event






08 / 12 / 2026

↳ Ship Your First Eval

Agent Code

```
1  // Inngest Function
2  export const agentRun = inngest.createFunction(
3    { id: "agent-run", triggers: { event: "agent/
4      run.requested" } },
5    async ({ event, step }) => {
6      // Extract event data
7      const { prompt, model } = event.data;
8
9
10     // Model called in step
11     const reply = await step.run("call-model",
12     async () => {
13       return await chat(prompt, model);
14     });
15     return { reply };
16   },
17 );
18
```

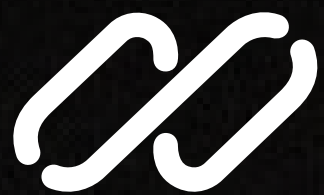


08 / 12 / 2026

↳ Ship Your First Eval

Scorer Code

```
1  // Inngest Scorer Function
2  export const feedbackScorer = createScorer(
3    inngest,
4    {
5      id: "feedback-scorer",
6      schema: z.object({ messageId: z.string() }),
7    },
8    async ({ event, step }) => {
9
10   // Waits for event
11     const feedback = await step.waitForEvent("wait-
12   for-feedback", {
13       event: "chat/feedback.clicked",
14       timeout: "1d",
15       if: `async.data.messageId ==
16   '${event.data.messageId}'`,
17     });
18
19     if (!feedback) return null;
20
21   // Return 0-1 Score
22     return {
23       name: "thumbs-up",
24       value: feedback.data.up ? 1 : 0,
25     };
26   },
27 );
```





08 / 12 / 2026

↳ Ship Your First Eval

Adding the Eval

```
1 // Inngest Function
2 export const agentRun = inngest.createFunction(
3   { id: "agent-run", triggers: { event: "agent/
4     run.requested" } },
5   async ({ event, step }) => {
6     // Extract event data
7     const { prompt, model } = event.data;
8
9
10    // Model called in step
11    const reply = await step.run("call-model",
12      async () => {
13        return await chat(prompt, model);
14      });
15    // Return response back to user
16    return { reply };
17  },
18 );
19
```

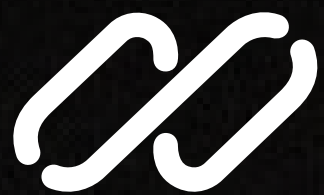


08 / 12 / 2026

↳ Ship Your First Eval

Adding the Eval

```
1 // Inngest Function
2 export const agentRun = inngest.createFunction(
3   { id: "agent-run", triggers: { event: "agent/
4     run.requested" } },
5   async ({ event, step }) => {
6     // Extract event data
7     const { prompt, model } = event.data;
8
9
10    // Model called in step
11    const reply = await step.run("call-model",
12      async () => {
13        return await chat(prompt, model);
14      });
15    // Defer scoring
16    defer("score-feedback", {
17      function: feedbackScorer,
18      data: { messageId: event.data.messageId },
19    });
20
21    // Return response back to user
22    return { reply };
23  },
24 );
25
```

08 / 12 / 2026

↳ Ship Your First Eval

Bonus

Model Split Experiment

```
1 // Inngest Function
2 export const recordMessage = inngest.createFunction(
3   { id: "record-message", triggers: { event: "chat/message.requested" } },
4   async ({ event, step, defer, group }) => {
5
6     // Set up the experiment
7     const { result: reply, variant, experimentRef } =
8       await group.experiment("chat-model", {
9
10      // Set up the model variants
11      variants: {
12        control: () => step.run("call-model-control", () =>
13          chat(event.data.prompt, DEFAULT_MODEL)),
14
15        candidate: () => step.run("call-model-candidate", () =>
16          chat(event.data.prompt, CANDIDATE_MODEL)),
17      },
18
19      // Split the Traffic
20      select: experiment.weighted({ control: 50, candidate: 50 }),
21    });
```



08 / 12 / 2026

↳ Ship Your First Eval

Viewing the Data

Pick First Eval


Adding Scoring Code

Viewing Eval Data



08 / 12 / 2026

↳ Ship Your First Eval



Runs by model

Model	Runs
vidia/nemotron...uper-120b-a12b	7
qwen/qwen3.8-max	4
vidia/nemotron...ltra-550b-a55b	1

Tokens by model

Model	Input	Output
vidia/nemotron...uper-120b-a12b	169	6,102
qwen/qwen3.8-max	253	1,596
vidia/nemotron...ltra-550b-a55b	26	1,024

AI Overview Panel

General information about your agent runs.



08 / 12 / 2026

↳ Ship Your First Eval

Workflow

Overview

Apps

Functions

Runs

Event Types

Events

Monitor

Metrics

Insights

AI

AI Overview

Experiments

Scores

Sessions

Sandboxes

Manage

Webhooks

Keys

All experiments > chat-model

chat-model

Last 7d All variants

Recommended: candidate

Total Run Count

27

TOTAL RUNS

candidate

16 runs

control

11 runs

thumbs-up

candidate

control

0

0.25

0.5

0.75

1

Score Summary

candidate

control

0

25

50

75

100

thumbs-up

variants

Adjust the scoring weight via the column title or sidebar.

Score	Variant	thumbs-up
94	candidate	0.94 Best User Satisfaction
91	control	0.91 Worst User Satisfaction

Open with insights



08 / 12 / 2026

↳ Ship Your First Eval

chat-model-v3

Last 7d

All variants


Recommended: control

Total Run Count




control 35 runs
candidate 34 runs


Score Summary




answered-question




clarity




conciseness




context-awareness




helpfulness




latency-ms



reply-words



thumbs-up

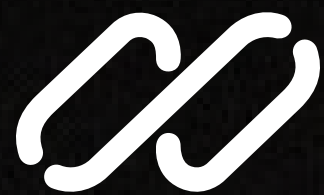


Variants

Adjust the scoring weight via the column title or sidebar.

Open with insights

Score	Variant	answered-question	clarity	conciseness	context-awareness	helpfulness	latency-ms	reply-words	thumbs-up
83	control	1 Best answered-question	0.97 Best clarity	0.95 Best conciseness	0.86 Worst context-awareness	1 Best helpfulness	1,509.91 Best latency-ms	24.06 Best reply-words	0.76 Worst thumbs-up
73	candidate	1 +0.0% vs best	0.9 Worst clarity	0.77 Worst conciseness	0.91 Best context-awareness	0.98 Worst helpfulness	2,838.85 Worst latency-ms	52.56 Worst reply-words	0.93 Best thumbs-up



08 / 12 / 2026

↳ Ship Your First Eval

RECAP

Decide on an Evaluation Metric

What measurements do you care about?

Create Scoring Function

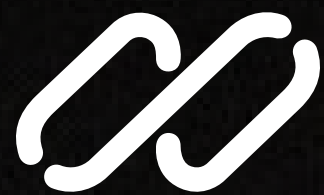
Wrap our scoring code.

Split Test Models in our Agent

Test out which model works best for our use case.

View Eval Results

View our eval data.



08 / 12 / 2026

↳ Ship Your First Eval

RECAP

Add More Metrics

Get a full picture of your agent performance by adding more data points.

Create Experiments

Experiment with new prompt versions and context.

Gather Scoring Data into a Dataset

Build your dataset of best responses and behaviors.

Use Findings for Improvement

Use your new dataset to fine-tune the model itself or modify the prompt and context.

INGEST

Thank you

Pick First Eval

Adding Scoring Code

Viewing Eval Data

